
Reinforcement Learning 2023, Master CS, Leiden University

Assignment 2 on Deep Q Learning (DQN)

Tom Stein (s3780120)¹ Tom Stein (s3780120)¹ Tom Stein (s3780120)¹

Abstract

This assignment report focuses on Deep Q Learning (DQN) with an application to the CartPole environment. The basic concept of DQN is introduced along with the experience replay and target network improvements. A hyperparameter scan is used to empirically compare the performance of the different models and discuss the results. A high degree of instability in the training process was observed, which is, to some extent, mitigated by specific adjustments of the parameters.

1. Introduction

In this assignment, an agent has to learn how to balance a pole in the vertical position. The environment (Brockman et al., 2016) presented is a well known and defined physics problem, being a reverse pendulum. The vertical position corresponds to the unstable equilibrium point and the pole itself is attached through a joint to a cart. The goal of this learning task is to keep upright the pole while moving the cart left or right. The possible action space is made up by a set of only two possible movements 0, 1, where 0: the cart is pushed to the left; 1: the cart is pushed to the right. The state space is composed by four values: position and velocity of the cart, angle and angular velocity of the pendulum. The agent receives a reward of +1 for every action performed resulting in keeping the pole into an equilibrium state. The environment is reset to the initial condition every time that the value of the angle between the pole and the vertical line is bigger than 15 or when the cart leaves the range $(-2, 4; 2, 4)$.

To tackle this problem tabular methods are not sufficient anymore, since keeping in memory all the possible states is not feasible anymore. Therefore, actions cannot be performed anymore on the basis of the Q-value stored in the table, but we need a way to generalize to unseen states. This can be achieved through the application of an Artificial Neu-

ral Network. In this assignment our deep neural network taken as input a state returns the Q values for the actions.

As a baseline comparison, we will use a random policy which has a reward around 22.

The structure of the following chapters is as follows. In section 2, 3 and 4, for each of the DQN algorithm hyperparameters optimization will be analyzed, highlighting the optimal configuration. Section 2 covers the basic case of DQN algorithm while, in section 3 the replay buffer is present and in section 4 the target network is considered. After all this architecture has been presented, section 5 is dedicated to the comparison and analysis of the different DQN approaches showing their pros and their limitations.

References

- Brockman, G., Cheung, V., Pettersson, L., Schneider, J., Schulman, J., Tang, J., and Zaremba, W. Openai gym, 2016.
- Plaat, A. *Deep Reinforcement Learning*. Springer, 2022. ISBN 978-981-19-0637-4. doi: 10.1007/978-981-19-0638-1. URL <https://doi.org/10.1007/978-981-19-0638-1>.

A. Hyperparameter Scan Results

^{*}Equal contribution ¹Faculty of Science, Leiden University, Leiden, The Netherlands. Correspondence to: Tom Stein <tom.stein@tu-dortmund.de>.